

CSI 402 – Lecture 11  
(Unix – Discussion on Files – continued)

# User and Group IDs

**Ref:** Chapter 3 of [HGS].

- Each user is given an ID (integer) called `uid`. (Most system programs use `uid` instead of the user's login name.)
- The special `uid 0` (zero) is given to the super user (`root`).
- Each user also belongs to at least one group. Each group has an ID (integer) called `gid`.
- When a file is created, the `uid` and the `gid` of the user who created the file are stored as part of the information regarding the file.

# More on File Modes

- **Recall:** Permissions (modes) can be specified as octal values (for user, group and other).
- Symbolic form also possible (constants defined in `<sys/stat.h>`):

<code>S_IRUSR</code>	<code>S_IWUSR</code>	<code>S_IXUSR</code>	<code>S_IRWXU</code>
<code>S_IRGRP</code>	<code>S_IWGRP</code>	<code>S_IXGRP</code>	<code>S_IRWXG</code>
<code>S_IROTH</code>	<code>S_IWOTH</code>	<code>S_IXOTH</code>	<code>S_IRWXO</code>

## Sample code segment:

```
int  fd;
mode_t fmode = S_IRWXU | S_IRGRP | S_IROTH;
if ((fd = open("file", O_WRONLY, fmode)) == -1)
    fprintf(stderr, "Error in open\n");  exit(1);
}
```

# Permission Bits for Executable Files

- Three extra permission bits, usually for executable files.
  - Set uid on execution: 04000      S\_ISUID
  - Set gid on execution: 02000      S\_ISGID
  - Save-text-image ("sticky bit"): 01000      S\_ISVTX

## Purpose of the Setuid Bit:

- If S\_ISUID permission for a file is set, the **effective uid** of the process executing the file is the uid of the owner of the file (rather than the uid of the user who started the execution).
- The process is given file access privileges of the file owner.

# Permission Bits for Executable Files (continued)

**Example:** The passwd program.

- Program used to change a user's password.
- Changing password requires permission to write to a file containing password information. (Normally, only root has permission to write to this file.)
- The executable file for the program is owned by root but has its S\_ISUID permission set.
- When a normal user runs the program, the program is given the effective user id of the owner (root). So, entries in the password file can be modified.

# Permission Bits for Executable Files (continued)

**Other examples:** The following programs also (generally) have their S\_ISUID permission set.

- Program to print files (usually `lpr`).
- Programs that update files containing top scorer information for some games.

**Setgid bit:** Similar to Setuid bit.

**Sticky Bit:**

- **Original purpose:** An executable with the sticky bit set is saved in the **swap area** of the disk (so that it can be loaded into memory quickly).
- With virtual memory, this bit is redundant for executable files.

# File Creation Mask (`umask`)

- When a file is created, permissions specified by mode are always modified using the 9-bit quantity called `umask`.
- Generally, value of `umask` is specified in a start up file such as `.bash_profile`.
- Value can be found using the shell command `umask`.  
(Usual default value: 022)
- Actual mode used for the file is given by

$$\text{given\_mode} \ \& \ \sim\text{umask}$$

- Setting `umask` to 022 prevents the open system call from creating a file that can be written by anyone other than the owner.

# File Creation Mask (continued)

- System call for `umask`:

```
mode_t  umask (mode_t  newmask)
```

- The call changes `umask` to the value specified by the parameter and returns the old value of `umask`.

**Reading Assignment:** Program example on page 45 of [HGS].



## Additional comments regarding open

- The call to `open` checks whether the specified access flags satisfy the file permissions.
- If not, the call returns `-1` and `errno` is set to `EACCES` (“permission denied” error).
- Special flag `O_EXCL` prevents creation of file if the file already exists.
- If flag is given as

`O_WRONLY | O_CREAT | O_EXCL`

and the file exists, `open` fails and `errno` is set to `EEXIST`.

## System call `chmod`:

### ■ **Prototype:**

```
int  chmod (const char *pathname,  mode_t  newmode)
```

- Can be used to change the permission bits of a file (including the special permission bits for executable files).
- The parameters represent the file name and the new permission bits respectively.
- The call succeeds only when issued by the owner of the file or the superuser (root).
- The value of `umask` is *not* used by `chmod`.
- Returns 0 if successful and -1 otherwise.

# Additional System Calls (continued)

## System call `chown`:

- **Prototype:**

```
int  chown (const char *pathname,  uid_t  newowner,  
           gid_t  newgroup)
```

- Can be used to change the owner and the group of a file.
- The parameters represent the file name, the uid of the new owner and the gid of the new group respectively.
- If the `newowner` parameter is `-1`, the owner is not changed. A similar comment applied to the `newgroup` parameter.

## System call `chown` (continued):

- **Caution:** Once the owner is changed, the change cannot be undone by the previous owner.
- The call succeeds only when issued by the owner of the file or the superuser (root); otherwise, `errno` is set to `EPERM`.
- Returns 0 if successful and -1 otherwise.

# Hard Links and Symbolic (Soft) Links

- **Link:** A mechanism that allows a file to be referred to by different names.
- **Note:** There is only one copy of the file; the links refer to the same copy.
- Useful in allowing different software teams to work together.
- Shell command for creating a hard link:

```
% ln oldfile secondname
```

- Shell command for creating a symbolic (or soft) link:

```
% ln -s oldfile anothername
```

- **Caution:** If a file is deleted, the symbolic links for the file become dangling pointers.

# Hard Links and Symbolic Links (continued)

## Hard Links and Symbolic links: Differences

- A normal user *cannot* create a hard link to a directory; however, a user may create a symbolic link to a directory.
- A hard link must be to a file within the same file system. A soft link may go across file systems.
- The number of hard links to a file is stored as part of the information regarding the file. The system does not store information about the number of symbolic links.

# Additional System Calls (continued)

## System call link:

### ■ **Prototype:**

```
int link (const char *original_path,  
          const char *new_path)
```

- Creates a hard link to an existing file.
- The parameters represent the name of an existing file and that of the link respectively.
- The link may be in a different directory.
- Returns 0 if successful and -1 otherwise.

## Some comments regarding unlink:

- The call to `unlink` only removes the link named in the call.
- The number of hard links to the file is decremented by 1.
- The disk space allocated to the file is reclaimed only if both of the following conditions hold:
  - 1 The hard link count has dropped to zero.
  - 2 No process has the file open.



# Additional System Calls (continued)

## System call symlink:

### ■ **Prototype:**

```
int  symlink (const char *real_name,  
              const char *sym_name)
```

- Creates a symbolic link to an existing file.
- The parameters represent the name of an existing file and that of the symbolic link respectively.
- The link may be in a different file system.
- Returns 0 if successful and -1 otherwise.

# Additional System Calls (continued)

## System calls stat, fstat and lstat:

### ■ Prototypes:

```
int  stat (const char *pathname, struct stat *buf)
int  fstat (int filedes,  struct stat *buf)
int  lstat (const char *pathname, struct stat *buf)
```

- Allow us to obtain information about files.
- In all three calls, the second parameter is a pointer to the buffer where information about the file can be stored. (The space for the buffer must have been allocated before the call.)
- stat: The first parameter represents the name of an existing file.
- fstat: The first parameter represents the file descriptor of an open file.

# Additional System Calls (continued)

## System calls stat, fstat and lstat (continued):

- lstat: The first parameter represents the name of a symbolic link. The call returns information about the link rather than the file pointed to by the link.
- All the calls return 0 if successful and -1 otherwise.

Data Members of struct stat: See Handout 11.1.  
(Additional discussion appears on page 54 of [HGS].)

Program example: Handout 11.2.

Reading Assignment: Program example on pages 56–57 of [HGS].