

CSI 402 – Lecture 12

(Unix – Directories and File Systems)

Ref: Chapter 4 of [HGS].

- Directories in Unix are also files. (However, they shouldn't be used like ordinary files.)
- Each directory entry stores the file name and the Inode number of the file.
- **Inode:**
 - Data structure that stores information about the file.
 - The `stat` system call returns the information from the Inode.

Directories (continued)

- Older versions of Unix:
 - File name restricted to 14 characters.
 - Directory entries are of fixed length.
- BSD 4.3 and later versions of Unix:
 - File names may have up to 256 characters.
 - Each directory entry has Inode number, the length of file name (1 byte) and the actual file name (string).
 - Directory entries are of variable length.
- **Caution:** Different versions of Unix may implement directories differently. It is best to use system calls to deal with directories.

Permission Bits for Directories:

- **Read permission:** Allows a user to list the files in the directory.
- **Write permission:** Allows a user to create new files and delete files in the directory.
- **Execute permission:** Allows a user to `cd` to the directory.

Sticky Bit and Directories:

- Why?
 - Some directories (e.g. /tmp) allow any user to create files. (The files are periodically removed by the root.)
 - The permission bits for such directories are read, write and execute for everyone.
 - However, a normal user should not be able to delete or rename files owned by others.
- If the sticky bit for a directory is set, then a file in the directory can be removed or renamed by a normal user only when the user owns the file.

Directory Related System Calls

- **Header file:** <dirent.h>
- Each directory entry is of type struct dirent with two data members:

```
ino_t  d_ino;    /* Inode number. */  
char  d_name[]; /* File name.     */
```

- If the value of d_ino is 0, then the entry does *not* correspond to a valid file.
- The string d_name is null terminated.

Directory Related System Calls (continued)

System call mkdir:

■ **Prototype:**

```
int  mkdir (const char *pathname,  mode_t mode)
```

- Creates a directory with name given by `pathname`.
- The permission bits for the created directory combine `mode` with `umask`.
- The created directory is initialized with two entries corresponding to `"."` and `".."`.

System call rmdir:

■ **Prototype:** `int rmdir (const char *pathname)`

- Removes the specified directory.
- A directory is removed only if it is empty (i.e., the only entries in the directory are for `"."` and `".."`).

Directory Related System Calls (continued)

System call opendir:

- **Prototype:** `DIR *opendir (const char *pathname)`
- Opens the specified directory.
- **Note:** Returns a pointer of type `DIR *`. The return value is `NULL` if an error occurs.

System call closedir:

- **Prototype:** `int closedir (DIR *dirptr)`
- Closes the directory specified by the the parameter.
- Returns 0 if successful and -1 otherwise.

Directory Related System Calls (continued)

System call readdir:

■ **Prototype:**

```
struct dirent * readdir (DIR *dirptr)
```

- Returns the next entry from the directory specified by the the parameter.
- **Note:** Returns a pointer of type struct dirent *. The return value is NULL if an error occurs or when there are no more entries in the directory.

System call rewinddir:

- **Prototype:** void rewinddir (DIR *dirptr)
- Goes back to the beginning of the directory specified by the parameter.
- The next call to readdir will return the first entry in the directory.

Program example: Handout 12.1.

Directory Related System Calls (continued)

System call chdir:

- **Prototype:** `int chdir (const char *path)`
- Changes the working directory to the one specified by the parameter `path`.
- Fails (and returns `-1`) if the parameter is not a valid directory or the user does not have execute permission for the directory.

System call getcwd:

- **Prototype:**

`char *getcwd (char *dname, size_t size);`

- Returns a pointer to the current directory path name; the path name is also copied into the array given by `dname`.
- Array `dname` should have size *at least one more than* the value of `size`.

Program example: Handout 12.2.

Directory Related System Calls (continued)

Function Pointers in C:

- Allow us to pass function names as parameters.
- **Example:** Handout 12.3.

System call `ftw`:

- **Header:** `<ftw.h>`
- **Prototype:**

```
int ftw (const char *path, int (*func)(), int depth)
```

- `ftw`: File tree walk.
- Performs a (recursive) traversal of the directory tree rooted at the path name given by `path`.

Directory Related System Calls (continued)

System call `ftw` continued:

- `depth`: Represents a limit on the number of file descriptors that `ftw` can use.
 - Value of 1 for `depth` will work, but may be too slow.
 - The value of `depth` can't be too large. (Each process may use only a limited number of file descriptors.)
- For each file visited during the traversal, the user defined function `func` will be called with three parameters:

```
int  func (const char *name, const struct stat *sptr,  
          int type)
```

Directory Related System Calls (continued)

System call `ftw` continued:

- `name`: Name of the file.
- `sptr`: Pointer to the `stat` structure for the file.
- `type`: Possible values are `FTW_F`, `FTW_D`, `FTW_DNR`, `FTW_SL` and `FTW_NS`. (See page 75 of [HGS] for the meanings of these constants.)
- The tree traversal continues if the user defined function returns the zero value; otherwise, `ftw` terminates the traversal.

Program Example: Handout 12.4.

Structure of an Inode

- Each file has an Inode which contains information about the file. (The `stat` system call obtains information about a file from the file's Inode.)
- **Size of each Inode:** 128 bytes (on most Unix systems).
- Each Inode contains
 - Information for the `stat` structure.
 - 12 **direct** pointers to data blocks.
 - One, two and three level **indirect** pointers to blocks.
- For each open file, the kernel keeps a copy of the corresponding Inode in memory.

Structure of Unix File System

- Disk divided into partitions; each partition is a “file system”.
- Each file system contains:
 - A boot block.
 - A super block which contains information about the state of the file system (e.g. the size of the file system, information regarding free blocks).
 - A sequence of Inodes.
 - A sequence of data blocks.