

CSI 402 – Lecture 15

(Unix – Processes and Pipes)

Ref: Chapter 7 of [HGS].

Pipe Mechanism in Unix:

- Allows a user to build complex commands by chaining together simple commands.
- Eliminates the need for temporary files.

Example 1: The following sequence of shell commands counts the number of lines containing the string "this" in the file "file.txt".

```
% grep this file.txt > temp
% wc -l temp
% rm temp
```

Processes and Pipes (continued)

Example 1 using pipes:

```
% grep  this  file.txt |  wc -l
```

Notes:

- **Pipe:** A first-in first-out buffer with a “read-end” and a “write-end”.
- To make the pipe in Example 1 to work:
 - The `stdout` for the `grep` program should be redirected to the write-end of pipe.
 - The `stdin` of the `wc` program should be redirected to the read-end of pipe.

General Information about Pipes

- There is *no* external file associated with a pipe; it is just a temporary internal buffer of limited capacity.
- The kernel blocks a process which tries to read from an empty pipe or which wants to write to a full pipe.
- The kernel also provides the synchronization needed among processes that use a pipe.
- A pipe can be used only between a process which created the pipe and the children of that process. (This restriction does *not* apply to “named pipes”.)

System call pipe

- **Prototype:** `int pipe (int filedes[2])`
- **Header:** `<unistd.h>`
- **filedes:** A two-element array of type `int`; the function fills in the two entries of this array.
- Each array element stores a file descriptor.
- `filedes[0]`: File descriptor for the read-end of the pipe (i.e., the descriptor to be used with the `read` system call).
- `filedes[1]`: File descriptor for the write-end of the pipe. (i.e., the descriptor to be used with the `write` system call).
- The function returns `-1` if the call fails.

System call pipe (continued)

Sample code segment:

```
int  pd[2];  char  x[10], y[10];

if (pipe(pd) == -1) {
    fprintf(stderr, "Failed pipe.\n"); exit(1);
}

strcpy(x, "012345678");

write(pd[1], x, 10);

read(pd[0], y, 10);

printf("%s\n", y);
```

System call `dup2`

- **Prototype:** `int dup2 (int fd1, int fd2)`
- **Header:** `<unistd.h>`.
- An existing file descriptor `fd1` is duplicated as file descriptor `fd2`.
- If the file (or device) corresponding to descriptor `fd2` is open, then it is closed.
- The duplication and closing are carried out as an **atomic** step.

Using dup2 for Redirection

(a) Redirecting stdin to a file:

```
int fd;
if ((fd = open("x.dat", O_RDONLY)) == NULL) {
    .
    . /* Code for error exit. */
}
if (dup2(fd, STDIN_FILENO) == -1) {
    .
    . /* Code for error exit. */
}
close(fd); /* Important! */
```

Note: At the end of the above segment (assuming that dup2 call was successful), all read operations to stdin will be redirected to the file "x.dat".

Using dup2 ... (continued)

(b) Redirecting stdout to a pipe:

```
int pd[2];
if (pipe(pd) == -1) {
    .
    . /* Code for error exit. */
}
if (dup2(pd[1], STDOUT_FILENO) == -1) {
    .
    . /* Code for error exit. */
}
close(pd[1]); /* Important! */
```

Note: At the end of the above segment (assuming that dup2 call was successful), all write operations to stdout will be redirected to the write-end of the pipe.

Program example: Handout 15.1.

Non-blocking Read/Write with Pipes

Reading Assignment: Program on pages 170–171 of [HGS].

- **Recall:** Read/write operation on a pipe may block the process.
- The `fcntl` system call can be used to make read/write non-blocking.

Sample Code Segment:

```
if (fcntl(pd[0], F_SETFL, O_NONBLOCK) == -1) {  
    .  
    .  /* Code for error exit. */  
}
```

Non-blocking Read/Write ... (continued)

Notes:

- **Non-blocking read from a pipe:** If pipe is empty, the read call returns immediately with value -1 and errno is set to EAGAIN.
- **Non-blocking write to a pipe:** If pipe is full, the write call returns immediately with value -1 and errno is set to EAGAIN.

Program Example: Handout 15.2.

Named Pipes (or FIFOs)

- Unnamed pipes can only be used between processes which have an ancestral relationship.
- Unnamed pipes are temporary; they need to be created every time and are destroyed when the corresponding processes exit.
- Named pipes (FIFOs) overcome both of these limitations.

Named Pipes (continued)

System call `mkfifo`:

- **Prototype:**

```
int mkfifo (const char *path, mode_t mode)
```

- **Headers:** `<sys/types.h>` and `<sys/stat.h>`.
- **path:** Name of the FIFO to be created.
- **mode:** File permissions for the FIFO; will be modified using `umask`.
- Returns `-1` if the call fails and `0` otherwise.

Named Pipes (continued)

Notes:

- The `open` system call is used to obtain a file descriptor for a FIFO.
- Reading data from a FIFO and writing data to a FIFO can be done using `read` and `write` system calls.
- Non-blocking read/write on a FIFO can be done by using `O_NONBLOCK` in the call to `open`.

Program Example: Handout 15.3.

Handling Multiple Pipes/FIFOs

- A server may communicate with two more clients through two or more separate pipes/FIFOs.
- At any time, data may be available from more than one pipe/FIFO.
- The server may want to wait until data is available from at least one of the pipes/FIFOs.
- **A more general situation:** A server may want to monitor a set of file descriptors.
- System call `select` handles such situations.

Handling Multiple Pipes/FIFOs (continued)

Representing File Descriptor Sets:

- Usual bit representation for sets (one bit per element).
- If a bit is 1 (i.e., the bit is ON), then the process is “interested” in that file descriptor; otherwise, the bit is 0 (OFF).

(a) Data Type `fd_set`:

- To be used for representing and manipulating a set of file descriptors.
- **Header:** `<sys/time.h>`

Handling Multiple Pipes/FIFOs (continued)

(b) Permitted Operations:

- Initialize all bits to 0.

```
void  FD_ZERO (fd_set *f);
```

- Set a specific bit to 1.

```
void  FD_SET (int fd, fd_set *f);
```

- Set a specific bit to 0.

```
void  FD_CLR (int fd, fd_set *f);
```

- Check whether a specific bit is ON.

```
int  FD_ISSET (int fd, fd_set *f);
```

Handling Multiple Pipes/FIFOs (continued)

Structure timeval:

■ Struct Definition:

```
struct timeval {  
    long tv_sec;      /* Seconds. */  
    long tv_usec;     /* Micro-seconds. */  
};
```

- Used in select system call to specify how long the process should wait.
- The micro-seconds part can be used to specify fraction of a second.

System call `select`

- **Prototype:**

```
int select (int nfds, fd_set *rfdset, fd_set *wfdset,  
            fd_set *efdset, struct timeval *tout)
```

- **Header:** `<sys/time.h>`

- `nfds`: The set of file descriptors to be monitored will be a subset of $\{0, 1, \dots, \text{nfds} - 1\}$.
- For `nfds`, the symbolic constant `FD_SETSIZE` defined in `<sys/time.h>` may also be used.
- `rfdset`: Pointer to the set of file descriptors which should be monitored for **reading**.
- `wfdset`: Pointer to the set of file descriptors which should be monitored for **writing**.
- `efdset`: Pointer to the set of file descriptors which should be monitored for **errors**.

System call `select` (continued)

- `tout`:
 - Pointer to the `timeval` structure. (Memory for the structure must have been allocated by the caller.)
 - Specifies the timeout time.
 - If this pointer is `NULL`, the call waits until an event of interest occurs for any of the file descriptors.
- The call to `select` returns 0 if timeout occurs and -1 on error; otherwise, returns the number of file descriptors for which events of interest occurred.
- **Important Note:** The bit masks pointed to by `rfd`s, `wfd`s and `efd`s are all modified by the call. (So, a user should keep a copy of the original masks.)

Program Example: Handout 15.4.